

GCDs: Data Frame Manipulation and Wrangling

[Code ▼](#)

- **1 Libraries for Data Frame Manipulation**
- **2 A Grammar for Data Wrangling**
- **3 Some Common Ways for Selecting Variables Using `dplyr`**
 - 3.1 *Select Variables Starting with or Ending with Certain Characters*
 - 3.2 *Selecting and Selecting Out Variables By/Between Index*
 - 3.3 *Selecting and Selecting Out Variables By or Between Character Name*
 - 3.4 *Selecting and Selecting Out Variables Characters in Their Names*
 - 3.5 *Selecting and Selecting Out Variables by Type*
- **4 Cleaning Data**
 - 4.1 *Removing duplicate rows using `distinct()`*
- **5 Filtering using `dplyr` and Understanding Filtering Operators**
- **6 Piping Multiple Filter Function Calls**
- **7 Filtering Cases by Character Names/String Values**
 - 7.1 *Filter Cases using `==`*
- **8 Filtering Cases by Value**
 - 8.1 *Filter by `<` and `>` or `<=` or `>=` ...*
- **9 Filter Cases by Conditional X or Y Using `|` Operator...**
 - 9.1 *Filter by range using `%in%` ...*
 - 9.2 *Filter by conditional X and Y using `&` operator...*
- **10 Filter Cases containing characters Using `str_detect()` and `%in%` ...**
- **11 Filtering Missing Data (NA 's)**
 - 11.1 *Filter using `na.omit()` ...*
 - 11.2 *Filter using `is.na()` and `!is.na()` ...*
 - 11.3 *Filter using `complete.cases()` ...*
- **12 Summarizing Data Using `dplyr`**
- **13 Summarizing `across()` Multiple Variables**
 - 13.1 *Summarize across by numeric variables...*
 - 13.2 *Summarize across by variable name vector...*
- **14 Summarizing Data using `group_by()`**
 - 14.1 *Passing functions in a `List()`*
 - 14.2 *Grouping by multiple variables`*
- **15 The Data Manipulation Workflow: Putting It All Together**

1 Libraries for Data Frame Manipulation

- `magrittr` : for piping data frame objects
- `dplyr` : for selecting, filtering, and mutating
- `stringr` : for working with strings

[Hide](#)

```
library(magrittr)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

Hide

```
library(stringr)
```

2 A Grammar for Data Wrangling

The `dplyr` package presents a type of grammar for wrangling data (H. Wickham and Francois 2020). `dplyr` is part of the `tidyverse` ecosystem and loads using `library(tidyverse)`. `dplyr` also imports functions from `tidyselect`, which is also part of the `tidyverse` ecosystem.

The functions from the packages in the `tidyverse` can be used without referencing the packages explicitly but to avoid functions of the same name from other packages being confused, we will reference the package/library and the function (e.g., `package::function()`).

Working with data involves creating new column variables, removing or renaming column variables, removing or deleting row observations, sorting, and summarizing data, often by groups. Consequently, there are five main function verbs for working with data in a data frame: `select`, `filter`, `mutate`, `arrange`, and `summarize`.

- `select(dataframe, variables_to_select)` : subset by columns
- `mutate(dataframe, variables_to_create)` and `dplyr::rename()` : add or modify existing columns
- `filter(dataframe, rows_to_select)` : subset by rows
- `arrange(dataframe, variable_to_sort_by)` : sort rows
- `summarize(dataframe, procedures_for_summarizing)` in conjunction with `dplyr::group_by()` : aggregate the data in some way

3 Some Common Ways for Selecting Variables Using `dplyr`

Using `select()`, you can select columns/variables from a data frame. The variables you select are retained and those that you don't select are not included in the returned data frame.

If not using `%>%`, the first argument passed into `select()` will be the data frame from which to select the

variables and the second and subsequent arguments can be variables.

```
select(mydataframe, myvars)
```

If piping a data frame with `magrittr`, the “.” or “.data” will serve to reference the inherited data frame.

```
dataframe %>% select(., myvars)
```

Variables can be passed separately without quotes or collectively as a character vector.

Hide

```
# passing variables by name (does not work with base R manipulation)
```

```
# passing variables separately
```

```
USArrests %>%  
  select(., Murder, Assault) %>% head()
```

```
##           Murder Assault  
## Alabama      13.2      236  
## Alaska       10.0      263  
## Arizona       8.1      294  
## Arkansas      8.8      190  
## California    9.0      276  
## Colorado      7.9      204
```

Hide

```
# passing variables separately as characters
```

```
# though combining them with c() is probably more clear
```

```
USArrests %>%  
  select(., "Murder", "Assault") %>% head()
```

```
##           Murder Assault  
## Alabama      13.2      236  
## Alaska       10.0      263  
## Arizona       8.1      294  
## Arkansas      8.8      190  
## California    9.0      276  
## Colorado      7.9      204
```

Hide

```
# passing a character vector
```

```
USArrests %>%  
  select(., c("Murder", "Assault")) %>% head()
```

```
##           Murder Assault
## Alabama    13.2    236
## Alaska     10.0    263
## Arizona     8.1    294
## Arkansas    8.8    190
## California  9.0    276
## Colorado   7.9    204
```

Hide

```
# passing an object holding a character vector
keep_vars <- c("Murder", "Assault") %>% head()
```

```
USArrests %>%
  select(., keep_vars)
```

```
## Note: Using an external vector in selections is ambiguous.
## i Use `all_of(keep_vars)` instead of `keep_vars` to silence this message.
## i See <https://tidyselect.r-lib.org/reference/faq-external-vector.html>.
## This message is displayed once per session.
```

##	Murder	Assault
## Alabama	13.2	236
## Alaska	10.0	263
## Arizona	8.1	294
## Arkansas	8.8	190
## California	9.0	276
## Colorado	7.9	204
## Connecticut	3.3	110
## Delaware	5.9	238
## Florida	15.4	335
## Georgia	17.4	211
## Hawaii	5.3	46
## Idaho	2.6	120
## Illinois	10.4	249
## Indiana	7.2	113
## Iowa	2.2	56
## Kansas	6.0	115
## Kentucky	9.7	109
## Louisiana	15.4	249
## Maine	2.1	83
## Maryland	11.3	300
## Massachusetts	4.4	149
## Michigan	12.1	255
## Minnesota	2.7	72
## Mississippi	16.1	259
## Missouri	9.0	178
## Montana	6.0	109
## Nebraska	4.3	102
## Nevada	12.2	252
## New Hampshire	2.1	57
## New Jersey	7.4	159
## New Mexico	11.4	285
## New York	11.1	254
## North Carolina	13.0	337
## North Dakota	0.8	45
## Ohio	7.3	120
## Oklahoma	6.6	151
## Oregon	4.9	159
## Pennsylvania	6.3	106
## Rhode Island	3.4	174
## South Carolina	14.4	279
## South Dakota	3.8	86
## Tennessee	13.2	188
## Texas	12.7	201
## Utah	3.2	120
## Vermont	2.2	48
## Virginia	8.5	156
## Washington	4.0	145
## West Virginia	5.7	81
## Wisconsin	2.6	53
## Wyoming	6.8	161

3.1 Select Variables Starting with or Ending with Certain Characters

One thing about `dplyr`, when you load the library, there are functions from other libraries that are imported along with `dplyr`'s own functions. These important functions are designed to work with each other, so the people who maintain the libraries have packaged them up nicely so you don't have to load separate libraries.

Many of the functions are imported from the `tidyselect` library and these functions are what give you additional manipulation ability. Some imported functions are: `all_of()`, `any_of()`, `contains()`, `ends_with()`, `everything()`, `last_col()`, `matches()`, `starts_with()`.

With functions like `starts_with()`, `contains()`, and `ends_with()`, you can select variables with patterns in their names.

Rather than passing the names of the variables as the second argument (e.g., `c("Murder", "Assault")`), you would pass the helper function, say `starts_with()`. Whatever `starts_with()` returns is what gets passed to `select()` as the variables. This is what is referred to as functional programming. Rather than coding specifically what to do, you will utilize the task of another function to pass its returned object as an argument to another function.

But first, we need to see what's going on with these functions, like `starts_with()`.

```
starts_with(match, ignore.case = TRUE, vars = NULL)
```

Notice the arguments we need to pass:

- `match`: A character vector
- `ignore.case`: If `TRUE`, the default, ignores case when matching names. This is most flexible.
- `vars`: A character vector of variable names. If not supplied, the variables are taken from the current selection context (as established by functions like `select()` or `pivot_longer()`).

Let's just try out `starts_with()` on its own. Let's set a required pattern `match = some character` and because `vars = NULL` by default, let's just set `vars = some character vector`. Note that `vars` is not the second argument, so you will want to name it in the function call.

Hide

```
starts_with(match = "a", vars = c("Hello", "Hi", "Bye"))
```

```
## integer(0)
```

Returns `integer(0)` which is speak for "there is no match". Hmm, OK. Let's try another character.

Hide

```
starts_with(match = "b", vars = c("Hello", "Hi", "Bye"))
```

```
## [1] 3
```

OK, so now an integer is returned (yes, try `is.integer()` if you don't believe me).

[Hide](#)

```
is.integer(starts_with("b", vars = c("Hello", "Hi", "Bye")))
```

```
## [1] TRUE
```

Importantly, the value refers to the element index/position in the `vars` vector. Because the third string "Bye" starts with "b", that's what is returned.

Try something else...

[Hide](#)

```
starts_with("h", vars = c("Hello", "Hi", "Bye"))
```

```
## [1] 1 2
```

Now a vector with length = 2 is returned, representing both the first and the second elements start with "h".

[Hide](#)

```
length(starts_with("h", vars = c("Hello", "Hi", "Bye")))
```

```
## [1] 2
```

See, it's really a vector containing the element(s) of the `vars` vector matching the pattern.

And yes, this the letter casing is ignored because the default `ignore.case = TRUE`. Set to `FALSE` if you want your match to be case sensitive.

[Hide](#)

```
starts_with("h",  
  vars = c("Hello", "Hi", "Bye"),  
  ignore.case = F)
```

```
## integer(0)
```

OK, no matches.

You will typically use `starts_with()` along with other functions. When using `starts_with()` in the context of `select()`, the `vars` argument is essentially passing
`vars = the names of the columns of the data frame passed to select()`.

Example:

```
select(mydataframe, starts_with(match = "my pattern", vars = "var names of mydataframe"))
```

Without piping...

[Hide](#)

```
select(USArrests, starts_with("m")) %>% head()
```

```
##           Murder
## Alabama      13.2
## Alaska       10.0
## Arizona        8.1
## Arkansas        8.8
## California      9.0
## Colorado        7.9
```

With piping...

Hide

```
USArrests %>%
  select(., starts_with("m")) %>% head()
```

```
##           Murder
## Alabama      13.2
## Alaska       10.0
## Arizona        8.1
## Arkansas        8.8
## California      9.0
## Colorado        7.9
```

Hide

```
USArrests %>%
  select(., ends_with("t")) %>% head()
```

```
##           Assault
## Alabama        236
## Alaska         263
## Arizona         294
## Arkansas        190
## California      276
## Colorado        204
```

3.2 *Selecting and Selecting Out Variables By/Between Index*

- `select(., 1,2)` : select first and second
- `select(., c(1,2))` : select first and second
- `select(., -c(1,2))` : select out first and second
- `select(., 1:2)` : select first through second

- `select(., c(1:2))` : select first through second
- `select(., -c(1:2))` : select out first through second

Recommendation: use options utilizing `c()` as this habit will be more versatile with base R functionality.

Let's make a data frame to work with first.

Hide

```
data <- data.frame(
  Id = c(100, 101, 102, 103, 104, 100, 105),
  Sex = c('male', 'female', 'Male', NA, 'man', "male", "neither"),
  Age = c(25, 33, 27, 40, 44, 25, 40),
  Renting = c("yes", NA, "yes", NA, "no", "yes", "yes")
)
```

Hide

```
data %>%
  select(., 1,2) # select this and that
```

```
##      Id      Sex
## 1 100    male
## 2 101  female
## 3 102    Male
## 4 103    <NA>
## 5 104     man
## 6 100    male
## 7 105  neither
```

Hide

```
data %>%
  select(., c(1,2)) # select this and that
```

```
##      Id      Sex
## 1 100    male
## 2 101  female
## 3 102    Male
## 4 103    <NA>
## 5 104     man
## 6 100    male
## 7 105  neither
```

Hide

```
data %>%
  select(., -c(1,2)) # select out this and that
```

```
##   Age Renting
## 1  25      yes
## 2  33    <NA>
## 3  27      yes
## 4  40    <NA>
## 5  44      no
## 6  25      yes
## 7  40      yes
```

Hide

```
data %>%
  select(., 1:2) # select from here to there
```

```
##   Id      Sex
## 1 100    male
## 2 101  female
## 3 102   Male
## 4 103    <NA>
## 5 104    man
## 6 100    male
## 7 105 neither
```

Hide

```
data %>%
  select(., c(1:3)) # select from here to there
```

```
##   Id      Sex Age
## 1 100    male  25
## 2 101  female  33
## 3 102   Male  27
## 4 103    <NA>  40
## 5 104    man  44
## 6 100    male  25
## 7 105 neither  40
```

Hide

```
data %>%
  select(., -c(1:3)) # select out from here to there
```

```
## Renting
## 1     yes
## 2    <NA>
## 3     yes
## 4    <NA>
## 5     no
## 6     yes
## 7     yes
```

3.3 Selecting and Selecting Out Variables By or Between Character Name

- `select(., "var1", "var2")`
- `select(., c("var1", "var2"))`
- `select(., -c("var1", "var2"))`
- `select(., var1:var2))`
- `select(., c("var1":"var2"))`
- `select(., -c("var1":"var2"))`

Recommendation: use options utilizing `c()` as this will be more versatile with base R functionality.

These also work but may lead to some confusion regarding usage of quotes:

- `select(., var1, var2)`
- `select(., c(var1, var2))`
- `select(., -c(var1, var2))`

Hide

```
data %>%
  select(., Id:Age) # select from here to there
```

```
##   Id    Sex Age
## 1 100  male  25
## 2 101 female  33
## 3 102  Male  27
## 4 103  <NA>  40
## 5 104   man  44
## 6 100  male  25
## 7 105 neither 40
```

Hide

```
data %>%
  select(., "Id":"Age") # select from here to there
```

```
##      Id      Sex Age
## 1 100    male  25
## 2 101  female  33
## 3 102   Male  27
## 4 103   <NA>  40
## 5 104    man  44
## 6 100    male  25
## 7 105 neither  40
```

Hide

```
data %>%
  select(., c("Id":"Age")) # select from here to there
```

```
##      Id      Sex Age
## 1 100    male  25
## 2 101  female  33
## 3 102   Male  27
## 4 103   <NA>  40
## 5 104    man  44
## 6 100    male  25
## 7 105 neither  40
```

Hide

```
data %>%
  select(., -c("Id":"Age")) # select out from here to there
```

```
##      Renting
## 1      yes
## 2   <NA>
## 3      yes
## 4   <NA>
## 5      no
## 6      yes
## 7      yes
```

Hide

```
# you can also use the ! operator to select NOT these variables (therefore, all others)
data %>%
  select(., !c("Id":"Age")) # select out from here to there
```

```
##   Renting
## 1     yes
## 2    <NA>
## 3     yes
## 4    <NA>
## 5     no
## 6     yes
## 7     yes
```

3.4 *Selecting and Selecting Out Variables Characters in Their Names*

- `select(., starts_with("character/s"))`
- `select(., ends_with("character/s"))`
- `select(., contains('e'))`

Hide

```
data %>% select(., starts_with('i'))
```

```
##   Id
## 1 100
## 2 101
## 3 102
## 4 103
## 5 104
## 6 100
## 7 105
```

Hide

```
data %>% select(., -starts_with('s'))
```

```
##   Id Age Renting
## 1 100  25     yes
## 2 101  33    <NA>
## 3 102  27     yes
## 4 103  40    <NA>
## 5 104  44     no
## 6 100  25     yes
## 7 105  40     yes
```

Hide

```
data %>% select(., ends_with('e'))
```

```
##   Age
## 1  25
## 2  33
## 3  27
## 4  40
## 5  44
## 6  25
## 7  40
```

Hide

```
data %>% select(., -ends_with('e'))
```

```
##   Id      Sex Renting
## 1 100    male     yes
## 2 101  female  <NA>
## 3 102    Male     yes
## 4 103   <NA>  <NA>
## 5 104     man     no
## 6 100    male     yes
## 7 105 neither     yes
```

Hide

```
data %>% select(., contains('g'))
```

```
##   Age Renting
## 1  25     yes
## 2  33  <NA>
## 3  27     yes
## 4  40  <NA>
## 5  44     no
## 6  25     yes
## 7  40     yes
```

Hide

```
data %>% select(., -contains('g'))
```

```
##   Id      Sex
## 1 100    male
## 2 101  female
## 3 102    Male
## 4 103   <NA>
## 5 104     man
## 6 100    male
## 7 105 neither
```

3.5 Selecting and Selecting Out Variables by Type

[Hide](#)

```
data %>% select(., where(is.numeric))
```

```
##      Id Age
## 1 100  25
## 2 101  33
## 3 102  27
## 4 103  40
## 5 104  44
## 6 100  25
## 7 105  40
```

[Hide](#)

```
data %>% select(., -where(is.numeric))
```

```
##      Sex Renting
## 1   male      yes
## 2 female    <NA>
## 3   Male      yes
## 4   <NA>    <NA>
## 5    man      no
## 6   male      yes
## 7 neither     yes
```

[Hide](#)

```
data %>% select(., where(is.character))
```

```
##      Sex Renting
## 1   male      yes
## 2 female    <NA>
## 3   Male      yes
## 4   <NA>    <NA>
## 5    man      no
## 6   male      yes
## 7 neither     yes
```

[Hide](#)

```
data %>% select(., -where(is.character))
```

```
##      Id Age
## 1 100  25
## 2 101  33
## 3 102  27
## 4 103  40
## 5 104  44
## 6 100  25
## 7 105  40
```

Hide

```
data %>% select(., where(is.logical))
```

```
## data frame with 0 columns and 7 rows
```

Hide

```
data %>% select(., -where(is.logical))
```

```
##      Id      Sex Age Renting
## 1 100    male  25      yes
## 2 101  female  33    <NA>
## 3 102    Male  27      yes
## 4 103    <NA>  40    <NA>
## 5 104     man  44      no
## 6 100    male  25      yes
## 7 105 neither  40      yes
```

4 Cleaning Data

Data files are messy and as a result require cleaning. You will have missing rows, incorrect variable names, files with columns named the same, NA s, strings for numbers, duplicate rows of data, people who completed a survey twice, and all sorts of unimaginable and unbelievable data problems. So cleaning is important.

Whereas `select()` is used for columns, `filter()` operates on rows. Data frame manipulation may involve keeping only certain rows for data, for example, male or female respondents, male respondents, those who do not contain missing values (e.g., NA s) for a specific column variable, who are of a certain age (or born in in certain year), who are above (or below) some acceptable criterion, etc.

When a column variable has more than one value (e.g., check using `unique()` to determine the unique elements contained), you may wish to filter on some but not others.

You may even need to filter rows in a data frame that are distinct (e.g., not duplicate responses). This is often a good first step in order to determine the size of the usable data set. `dplyr::distinct()` makes de-duplicating easy as this function will return only distinct rows.

4.1 *Removing duplicate rows using `distinct()`*

- `dplyr::distinct()` : remove duplicate rows
- `dplyr::distinct(., column)` : remove duplicate rows by column
- `na.omit()` : remove any row with NA's

Let's use the simple data data frame.

Hide

```
data %>% view(.) # the full data frame
```

Show entries

Search:

	Id	Sex	Age	Renting
	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>
1	100	male	25	yes
2	101	female	33	
3	102	Male	27	yes
4	103		40	
5	104	man	44	no
6	100	male	25	yes
7	105	neither	40	yes

Showing 1 to 7 of 7 entries

Previous

Next

Notice that rows 1 and 6 are the same person (e.g., Id) and have exactly the same data for all variables.

Hide

```
data[1,] == data[6,]
```

```
##      Id Sex Age Renting
## 1 TRUE TRUE TRUE    TRUE
```

Great that they are consistent but you don't want their data twice. So let's just remove any rows that are identical.

Hide

```
data %>%
  distinct(.) %>% # Remove exact duplicates
  view(.)
```

Show entries

Search:

	Id	Sex	Age	Renting
	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>
1	100	male	25	yes
2	101	female	33	
3	102	Male	27	yes
4	103		40	
5	104	man	44	no
6	105	neither	40	yes

Showing 1 to 6 of 6 entries

Previous Next

If you know each row is unique based on a variable in the data frame, you can use `distinct()` for that variable.

Hide

```
data %>%
  distinct(., Id) %>% view(.) # Remove duplicates by variable; passes unique values for data
frame
```

Show entries

Search:

	Id
	<input type="text" value="All"/>
1	100
2	101
3	102
4	103
5	104
6	105

Showing 1 to 6 of 6 entries

But this just returns the unique values in `Id`. To retain the variables, set `.keep_all = T`.

[Hide](#)

```
data %>%
  distinct(., Id, .keep_all = T) %>% view(.)
```

Show **100** entries

Search:

	Id	Sex	Age	Renting
	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>
1	100	male	25	yes
2	101	female	33	
3	102	Male	27	yes
4	103		40	
5	104	man	44	no
6	105	neither	40	yes

Showing 1 to 6 of 6 entries

[Previous](#)
[1](#)
[Next](#)

Notice, however, this only removed the last instance or `Id == 100`. Which row to include is a judgment call. The first, the last, neither, the average? Is there a correct answer?

5 Filtering using dplyr and Understanding Filtering Operators

Filtering cases using the `dplyr::filter()` verbs works by removing rows that do not match a specific criterion and then by returning the data frame that omits the mismatched condition.

Some useful filtering operators and functions include: `==`, `>`, `>=`, `&`, `|`, `!`, `xor()`, `c()`, `is.na()`, `between()`, `near()`.

Row/Observations/Cases can be filtered to “include” only certain matched conditions or can be filtered to “exclude” by negating those matched conditions. If the column variable `Sex` is in the data frame and cases are 'male', 'men', 'female', 'women', 'neither', NA, etc., you can specify the column `Sex` variable and then the row matching condition(s).

The first argument in `dplyr::filter()` is a data frame, and the function all

`dplyr::filter(data, Sex == 'female')` will filter the data frame named `data` to include rows for which the `sex` column equals `'female'`. In other words, `TRUE` rows.

Hide

```
dplyr::filter(data, Sex == 'female')
```

```
##      Id      Sex Age Renting
## 1 101 female  33    <NA>
```

Similarly, the function call `dplyr::filter(., Sex == 'male')` can be read “filter the data frame to include rows for which the value of `Sex == 'male'` is `TRUE`”.

More flexibly, however, you could specify a vector containing acceptable strings using `c()`.

`dplyr::filter(., Sex %in% c('male'))` filters the rows to include only those for which the value for `sex` is in the string vector which includes a single string, `'male'` whereas

`dplyr::filter(., Sex %in% c('male', 'Man'))` filters the rows to include only those for which the value for `Sex` is in the string vector which includes `'male'` and `'Man'`. Cases containing `'Male'`, `'Men'` (R is a case-sensitive language), or `'female'`, for example, will not be included in the returned data frame because they do not match values in the string vector.

6 Piping Multiple Filter Function Calls

In many cases, data filtering will involve different conditions for different column variables, so specifying them separately as separate lines of code is most appropriate.

When passing a data frame using `%>%` from `magrittr`, the first argument for the data frame can be specified using a `.` because the function inherits the data frame manipulated. However, `dplyr` also understand this so the `.` can also be omitted for convenience; this is the general practice you will see in forums like stackoverflow.com[stackoverflow.com]. Example to follow.

7 Filtering Cases by Character Names/String Values

7.1 Filter Cases using `==`

Hide

```
data %>%
  dplyr::filter(., Sex == 'female')
```

```
##      Id      Sex Age Renting
## 1 101 female  33    <NA>
```

Hide

```
# not equal
data %>%
  dplyr::filter(., Sex != 'female')
```

```
##      Id      Sex Age Renting
## 1 100    male  25     yes
## 2 102   Male  27     yes
## 3 104    man  44     no
## 4 100    male  25     yes
## 5 105 neither  40     yes
```

Hide

```
# multiple filters
data %>%
  dplyr::filter(., Sex == 'female' & Age > 27) # this "AND" that
```

```
##      Id      Sex Age Renting
## 1 101 female  33    <NA>
```

Hide

```
data %>%
  dplyr::filter(., Sex == 'female' | Age > 27) # this "OR" that
```

```
##      Id      Sex Age Renting
## 1 101 female  33    <NA>
## 2 103    <NA>  40    <NA>
## 3 104    man  44     no
## 4 105 neither  40     yes
```

A cleaner method involves separate lines of code. Although cleaner, this will not allow the “OR” option because the data frame that is returned from the first `filter()` is passed to the second `filter()` and all cases other than “female” have already been removed from the data frame.

Hide

```
data %>%
  dplyr::filter(., Sex == 'female') %>% # keep female (and add another pipe)
  dplyr::filter(., Age >= 27)           # keep only those equal to or older than 27
```

```
##      Id      Sex Age Renting
## 1 101 female  33    <NA>
```

8 Filtering Cases by Value

8.1 Filter by < and > or <= or >= ...

Hide

```
data %>% dplyr::filter(., Age < 40) # keep those Less than
```

```
##      Id      Sex Age Renting
## 1 100    male  25     yes
## 2 101 female  33    <NA>
## 3 102   Male  27     yes
## 4 100    male  25     yes
```

Hide

```
data %>% dplyr::filter(., Age > 40) # keep older than
```

```
##      Id Sex Age Renting
## 1 104 man  44      no
```

Hide

```
data %>% dplyr::filter(., Age >= 40) # keep equal to or older than
```

```
##      Id      Sex Age Renting
## 1 103    <NA>  40    <NA>
## 2 104     man  44      no
## 3 105 neither  40     yes
```

9 Filter Cases by Conditional X or Y Using | Operator...

Using the “OR” operator, | , cases can be included if “this” OR “that” condition.

Hide

```
data %>%
  dplyr::filter(., Age == 25 | Age == 40) # filter out numeric values IN a range
```

```
##      Id      Sex Age Renting
## 1 100    male  25     yes
## 2 103    <NA>  40    <NA>
## 3 100    male  25     yes
## 4 105 neither  40     yes
```

Hide

```
data %>%
  dplyr::filter(., Sex == 'male' | Sex == 'female')
```

```
##      Id    Sex Age Renting
## 1 100   male  25     yes
## 2 101 female  33    <NA>
## 3 100   male  25     yes
```

Hide

although dplyr::filter(sex %in% c('male', 'female')) would be easier

```
data %>%
  dplyr::filter(., Sex == 'male' | Age == 27)
```

```
##      Id    Sex Age Renting
## 1 100   male  25     yes
## 2 102   Male  27     yes
## 3 100   male  25     yes
```

Hide

```
data %>%
  dplyr::filter(., between(Age, 27, 33))
```

```
##      Id    Sex Age Renting
## 1 101 female  33    <NA>
## 2 102   Male  27     yes
```

Hide

```
data %>%
  dplyr::filter(., between(Age, 27, 33))
```

```
##      Id    Sex Age Renting
## 1 101 female  33    <NA>
## 2 102   Male  27     yes
```

9.1 Filter by range using %in% ...

Though less flexible than using `between()` ...

Hide

```
data %>%
  dplyr::filter(., Age %in% 20:43)    # filter out numeric values IN a range
```

```
##      Id      Sex Age Renting
## 1 100    male  25    yes
## 2 101  female  33    <NA>
## 3 102    Male  27    yes
## 4 103    <NA>  40    <NA>
## 5 100    male  25    yes
## 6 105 neither  40    yes
```

If a vector object is already defined (e.g., `my_levels = c('male', 'female')`), it can be used for filtering also. Such approaches are useful when data manipulation involves reusing a reference as it simplifies coding and reduces errors because the specification is defined only once.

Hide

```
my_levels = c('male', 'female')

data %>%
  dplyr::filter(., Sex %in% my_levels)
```

```
##      Id      Sex Age Renting
## 1 100    male  25    yes
## 2 101  female  33    <NA>
## 3 100    male  25    yes
```

When inclusion is inappropriate, exclusion may be useful. The `!` operator means “NOT” in R so it is great to accomplish the opposite. For example, `dplyr::filter(., !sex %in% c('male', NA))` will “filter the data frame to include rows in the `sex` column for which the value is NOT in the vector”.

Hide

```
# exclusion
data %>%
  dplyr::filter(., !Sex %in% c('male', NA)) # keep only if NOT in vector
```

```
##      Id      Sex Age Renting
## 1 101  female  33    <NA>
## 2 102    Male  27    yes
## 3 104    man  44    no
## 4 105 neither  40    yes
```

Hide

```
data %>%
  dplyr::filter(., !Sex %in% c('male', 'Men')) # keep only if NOT in vector
```



```
##      Id      Sex Age Renting
## 1 101  female  33    <NA>
## 2 102   Male   27     yes
## 3 103   <NA>  40    <NA>
## 4 104    man  44     no
## 5 105 neither  40     yes
```

9.2 Filter by conditional X and Y using & operator...

Hide

```
data %>%
  dplyr::filter(., Id >= 102 & Age <= 43)    # filter by range
```

```
##      Id      Sex Age Renting
## 1 102   Male   27     yes
## 2 103   <NA>  40    <NA>
## 3 105 neither  40     yes
```

Hide

```
data %>%
  dplyr::filter(., Age >= 20 & Age <= 43)    # filter by range
```

```
##      Id      Sex Age Renting
## 1 100   male   25     yes
## 2 101  female  33    <NA>
## 3 102   Male   27     yes
## 4 103   <NA>  40    <NA>
## 5 100   male   25     yes
## 6 105 neither  40     yes
```

Hide

#Note: Age 20:43 won't work. Can you figure out why?

10 Filter Cases containing characters Using str_detect() and %in%...

- stringr::str_detect() : returns matching conditions

Hide

```
data %>% dplyr::filter(., stringr::str_detect(Sex, "ma"))
```

```
##      Id      Sex Age Renting
## 1 100    male  25      yes
## 2 101 female  33    <NA>
## 3 104    man  44      no
## 4 100    male  25      yes
```

But the case for which `Sex = Male` is now missing. This is because `stringr::str_detect()` is a case-sensitive pattern match.

You can fix this in a couple ways:

1. make the cases in `Sex` all lower case to `mutate()` the fix or
2. to wrap `Sex` in `tolower()` to make cases lowercase.

The first option might be better if you want to fix the problem in the data frame.

Other casing functions are:

- `tolower()` : returns lower case of string
- `toupper()` : returns upper case of string
- `tools::toTitleCase()` : returns Title Case (capitalize first letter of string)

Hide

```
data %>%
  mutate(., Sex = tolower(Sex)) %>%
  filter(., stringr::str_detect(tolower(Sex), "ma"))
```

```
##      Id      Sex Age Renting
## 1 100    male  25      yes
## 2 101 female  33    <NA>
## 3 102    male  27      yes
## 4 104    man  44      no
## 5 100    male  25      yes
```

Hide

```
data %>%
  filter(., stringr::str_detect(tolower(Sex), "ma"))
```

```
##      Id      Sex Age Renting
## 1 100    male  25      yes
## 2 101 female  33    <NA>
## 3 102    Male  27      yes
## 4 104    man  44      no
## 5 100    male  25      yes
```

Hide

```
# along with toTitleCase() makes it better
data %>%
  mutate(., Sex = tools::toTitleCase(tolower(Sex))) %>%
  filter(., stringr::str_detect(tolower(Sex), "ma"))
```

```
##      Id      Sex Age Renting
## 1 100    Male  25      yes
## 2 101 Female  33     <NA>
## 3 102    Male  27      yes
## 4 104    Man  44      no
## 5 100    Male  25      yes
```

But notice the `male` and `man` issue is still a problem. Hard code a fix using `%in%` and `case_when()` ...

Hide

```
data %>%
  mutate(.,
    Sex = case_when(
      tolower(Sex) %in% c("male", "man") ~ "Male",
      tolower(Sex) %in% c("female", "woman") ~ "Female"
    ))
```

```
##      Id      Sex Age Renting
## 1 100    Male  25      yes
## 2 101 Female  33     <NA>
## 3 102    Male  27      yes
## 4 103  <NA>  40     <NA>
## 5 104    Male  44      no
## 6 100    Male  25      yes
## 7 105  <NA>  40      yes
```

Or for a more flexible fix using `str_detect()` ...

BUT beware that the order of operations matters here. Because the string “female” contains “ma”, you could accidentally recode all “female” cases to “male” if you perform the `case_when()` conversion on “male” first.

Hide

```
data %>%
  mutate(.,
    Sex = case_when(
      stringr::str_detect(tolower(Sex), "fe") ~ "Female",
      stringr::str_detect(tolower(Sex), "ma") ~ "Male",
    ))
```

```
##      Id    Sex Age Renting
## 1 100   Male  25     yes
## 2 101 Female  33    <NA>
## 3 102   Male  27     yes
## 4 103    <NA> 40    <NA>
## 5 104   Male  44     no
## 6 100   Male  25     yes
## 7 105    <NA> 40     yes
```

11 Filtering Missing Data (NA's)

`is.na()` will return a logical vector for which `TRUE` represents there are missing values.

Try on the entire data frame...

Hide

```
is.na(data)
```

```
##      Id    Sex Age Renting
## [1,] FALSE FALSE FALSE  FALSE
## [2,] FALSE FALSE FALSE   TRUE
## [3,] FALSE FALSE FALSE  FALSE
## [4,] FALSE  TRUE FALSE   TRUE
## [5,] FALSE FALSE FALSE  FALSE
## [6,] FALSE FALSE FALSE  FALSE
## [7,] FALSE FALSE FALSE  FALSE
```

You can see that some columns contain cases/rows with `TRUE` indicating the cell contains `NA`.

The negation operator, `!`, will be used to illustrate some filtering approaches. Because `filter()` will filter out `FALSE` cases and retain `TRUE` ones, so you may sometimes need to negate a function so that you keep the rows you want to keep.

- `na.omit()` : removes rows with NAs
- `dplyr::filter(., is.na(column_name))` : keep rows with NA in specific variable
- `dplyr::filter(., !is.na(column_name))` : remove rows with NA in specific variable
- `dplyr::filter(., complete.cases(.))` : remove rows

11.1 Filter using `na.omit()` ...

Hide

```
data %>%
  na.omit(.) %>%      # omit any rows with NAs
  view(.)
```

Show entries

Search:

	Id	Sex	Age	Renting
	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>
1	100	male	25	yes
3	102	Male	27	yes
5	104	man	44	no
6	100	male	25	yes
7	105	neither	40	yes

Showing 1 to 5 of 5 entries

Previous Next

11.2 Filter using `is.na()` and `!is.na()` ...

Hide

```
data %>%
  filter(., is.na(Sex)) %>%      # keep NAs by variable
  view(.)
```

Show entries

Search:

	Id	Sex	Age	Renting
	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>
1	103		40	

Showing 1 to 1 of 1 entries

Previous Next

Hide

```
data %>%
  filter(., !is.na(Sex)) %>%    # remove NAs by variable
  view(.)
```

Show entries

Search:

	Id	Sex	Age	Renting
--	----	-----	-----	---------

<div>All</div>		<div>All</div>		<div>All</div>		<div>All</div>	
1	100	male		25	yes		
2	101	female		33			
3	102	Male		27	yes		
4	104	man		44	no		
5	100	male		25	yes		
6	105	neither		40	yes		
Showing 1 to 6 of 6 entries				Previous	1	Next	

Hide

```
data %>%  
  filter(., !is.na(Sex)) %>%  
  filter(., !is.na(Renting)) %>%  
  view(.)
```

Show

100 ▾

 entries

Search:

Id		Sex	Age		Renting	
<div>All</div>		<div>All</div>	<div>All</div>		<div>All</div>	
1	100	male	25	yes		
2	102	Male	27	yes		
3	104	man	44	no		
4	100	male	25	yes		
5	105	neither	40	yes		
Showing 1 to 5 of 5 entries				Previous	1	Next

Hide

```
# separate calls on separate lines can also make code inclusion/exclusion easy
data %>%
  #filter(., !is.na(Sex)) %>%
  filter(., !is.na(Renting)) %>%
  view(.)
```

Show entries

Search:

	Id	Sex	Age	Renting
	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>	<input type="text" value="All"/>
1	100	male	25	yes
2	102	Male	27	yes
3	104	man	44	no
4	100	male	25	yes
5	105	neither	40	yes

Showing 1 to 5 of 5 entries

Previous

Next

11.3 Filter using `complete.cases()` ...

The `complete.cases()` function returns a logical vector for which `TRUE` reflects the row has complete information and no missing cases. Using `complete.cases()` along with `filter()`, you would retain all rows `TRUE` rows.

Hide

```
data %>%
  dplyr::filter(., complete.cases(.))
```

```
##   Id    Sex Age Renting
## 1 100   male  25     yes
## 2 102   Male  27     yes
## 3 104    man  44      no
## 4 100   male  25     yes
## 5 105 neither  40     yes
```

12 Summarizing Data Using `dp1yr`

If you want a quick summary of data, `summary()` will provide some basic information for you.

Hide

```
summary(data)
```

```
##           Id           Sex           Age           Renting
##  Min.      :100.0   Length:7      Min.      :25.00   Length:7
##  1st Qu.:100.5   Class :character  1st Qu.:26.00   Class :character
##  Median :102.0   Mode  :character  Median :33.00   Mode  :character
##  Mean    :102.1                                Mean    :33.43
##  3rd Qu.:103.5                                3rd Qu.:40.00
##  Max.    :105.0                                Max.    :44.00
```

However, there are many other ways to summarize data. To introduce data summary techniques using `dplyr`, we will open the `diamonds` data set from the `ggplot2` library. Then, we will use `dplyr::summarise()` or `dplyr::summarize()` to summarize the data. The `summarise()` function works similar to `mutate()` insofar as variables are created but differs insofar as the data frame returned from `summarise()` contains only the variable(s) referenced in `summarize()`.

In the example below, we summarize by creating a new variable which is set to represent some data summary technique. In essence, summarizing is for descriptive statistics. Using `mean()`, we can summarize the data by taking the mean of the `price` variable.

[Hide](#)

```
diamonds <- ggplot2::diamonds  # assign data to object

diamonds %>%
  summarise(.,
    mean = mean(price, na.rm = T),
  )
```

```
## # A tibble: 1 × 1
##   mean
##   <dbl>
## 1 3933.
```

Notice what is returned is a single value reflecting the mean of all the data in the data frame. We could have obtained the same without using `dplyr`.

[Hide](#)

```
mean(diamonds$price, na.rm = T)      # $ notation
```

```
## [1] 3932.8
```

But we lose flexibility of easily adding new summary procedures.

[Hide](#)


```
diamonds %>%
  summarise(.,
    mean = mean(price, na.rm = T),
    sd    = sd(price, na.rm = T)
  )
```

```
## # A tibble: 1 × 2
##   mean    sd
##   <dbl> <dbl>
## 1 3933. 3989.
```

Now there is a mean and standard deviation for price. You can also add the sample size using `dplyr::n()`.

Hide

```
diamonds %>%
  summarise(.,
    mean = mean(price, na.rm = T),
    sd    = sd(price, na.rm = T),
    n     = n()
  )
```

```
## # A tibble: 1 × 3
##   mean    sd    n
##   <dbl> <dbl> <int>
## 1 3933. 3989. 53940
```

13 Summarizing across() Multiple Variables

Summarizing a single variable is useful but if you want to summarize by many, you likely don't want to code a new line for each variable. In such cases, you can use `across()` as a helper function as was used for creating new variables with `mutate()` (see previous lesson).

Remember `across()` will want you to pass the columns to summarize by, `.cols`, the function for how to summarize, `.fns`, and the names for how to name the new variables, `.names` (which will be `NULL` by default). The `.x` here stands for passing the vector to the mean function and not the data frame. More on `~` and `.x` later.

13.1 Summarize across by numeric variables...

Hide

```
diamonds %>%
  summarise(., across(.cols = where(is.numeric),
    .fns = ~mean(.x, na.rm = TRUE))
  )
```

```
## # A tibble: 1 × 7
##   carat depth table price      x      y      z
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.798  61.7  57.5 3933.  5.73  5.73  3.54
```

Passing `.names` ...

Hide

```
diamonds %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
                        .fns = ~mean(.x, na.rm = TRUE),
                        .names = "{.col}_mean")
  )
```

```
## # A tibble: 1 × 3
##   carat_mean depth_mean table_mean
##   <dbl>      <dbl>      <dbl>
## 1    0.798      61.7      57.5
```

13.2 Summarize across by variable name vector...

Hide

```
diamonds %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
                        .fns = ~mean(.x, na.rm = TRUE),
                        .names = "{.col}_mean")
  )
```

```
## # A tibble: 1 × 3
##   carat_mean depth_mean table_mean
##   <dbl>      <dbl>      <dbl>
## 1    0.798      61.7      57.5
```

14 Summarizing Data using `group_by()`

A summary of the entire data set is fine for understanding grand mean and other metrics but this lacks information at a group level, for example, by sex, ethnicity, personality, city, job title, or in the `diamonds` data by `diamond cut`, `clarity`, or other variation.

When summarizing data, you will often want to summarize by levels of other variables, either categorical or numeric. In such cases, you can `group_by()` another variable and then summarize.

Let's summarize in several ways after grouping by `diamond cut`.

Hide

```
diamonds %>%
  group_by(., cut) %>%
  summarise(.,
    n = n(), # or also length()
    mean = mean(price, na.rm = T),
    sd = sd(price, na.rm = T)
  )
```

```
## # A tibble: 5 × 4
##   cut      n mean   sd
##   <ord>   <int> <dbl> <dbl>
## 1 Fair    1610 4359. 3560.
## 2 Good    4906 3929. 3682.
## 3 Very Good 12082 3982. 3936.
## 4 Premium 13791 4584. 4349.
## 5 Ideal   21551 3458. 3808.
```

You see that price increases with `cut` quality. Using `str()` or `glimpse()` you can see `cut` is an ordered factor.

Hide

```
diamonds %>% str()
```

```
## tibble [53,940 × 10] (S3: tbl_df/tbl/data.frame)
## $ carat : num [1:53940] 0.23 0.21 0.23 0.29 0.31 0.24 0.24 0.26 0.22 0.23 ...
## $ cut : Ord.factor w/ 5 levels "Fair"<"Good"<...: 5 4 2 4 2 3 3 1 3 ...
## $ color : Ord.factor w/ 7 levels "D"<"E"<"F"<"G"<...: 2 2 2 6 7 7 6 5 2 5 ...
## $ clarity: Ord.factor w/ 8 levels "I1"<"SI2"<"SI1"<...: 2 3 5 4 2 6 7 3 4 5 ...
## $ depth : num [1:53940] 61.5 59.8 56.9 62.4 63.3 62.8 62.3 61.9 65.1 59.4 ...
## $ table : num [1:53940] 55 61 65 58 58 57 57 55 61 61 ...
## $ price : int [1:53940] 326 326 327 334 335 336 336 337 337 338 ...
## $ x : num [1:53940] 3.95 3.89 4.05 4.2 4.34 3.94 3.95 4.07 3.87 4 ...
## $ y : num [1:53940] 3.98 3.84 4.07 4.23 4.35 3.96 3.98 4.11 3.78 4.05 ...
## $ z : num [1:53940] 2.43 2.31 2.31 2.63 2.75 2.48 2.47 2.53 2.49 2.39 ...
```

But wow, that was a lot of coding lines. You might want to write your code to be more flexible so every variable is not hard coded. Could you imagine having 50 variables?

14.1 *Passing functions in a List()*

Doing this requires a little fancy coding by passing two functions as a `list` object. A `list` is a special object (e.g., container) for which its elements can be different types of objects. Whereas elements of `vectors` can be only character or only numeric, elements of lists can hold different object. One element can be a numeric vector, another element a data frame, another element a character vector, etc. Many functions used in R will actually return lists for which elements contain different types of objects.

Back to two or more functions. If you pass a `list()` with arguments for the `mean` and the `sd`, you can summarize by both. if you want to prevent errors (yes you do) and want to keep the summaries separate, you

can modify `.names` to pass the column and the function, `"{.col}_{.fn}"`. The underscore is not needed; it only helps with readability of the variables. Let's pass the summary procedures as a `list`. Unfortunately, `n()` will throw an error but `length()` will also return the length of the vector, after of course, the grouping variable.

Here are a couple ways to do this. If you pass only the functions into the list, `.names` will take on the order of the functions such that mean would be `"_1"`. This is confusing.

Hide

```
diamonds %>%
  group_by(., cut) %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
                        .fns = list(mean, sd, length),
                        .names = "{.col}_{.fn}"))
)
```

```
## # A tibble: 5 × 10
##   cut    carat_1 carat_2 carat_3 depth_1 depth_2 depth_3 table_1 table_2 table_3
##   <ord>    <dbl>  <dbl>  <int>  <dbl>  <dbl>  <int>  <dbl>  <dbl>  <int>
## 1 Fair      1.05    0.516   1610   64.0    3.64   1610    59.1    3.95   1610
## 2 Good      0.849    0.454   4906   62.4    2.17   4906    58.7    2.85   4906
## 3 Very ...  0.806    0.459  12082   61.8    1.38  12082    58.0    2.12  12082
## 4 Premi...  0.892    0.515  13791   61.3    1.16  13791    58.7    1.48  13791
## 5 Ideal     0.703    0.433  21551   61.7    0.719  21551    56.0    1.25  21551
```

A better approach would be to assign the function a name in the `list()` function call so that the name is appended and the variable is named meaningfully.

Hide

```
diamonds %>%
  group_by(., cut) %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
                        .fns = list(mean = mean,
                                    sd = sd,
                                    n = length
                                    ),
                        .names = "{.col}_{.fn}"))
)
```

```
## # A tibble: 5 × 10
##   cut    carat...1 carat...2 carat_n depth...3 depth...4 depth_n table...5 table...6 table_n
##   <ord>    <dbl>    <dbl>  <int>  <dbl>    <dbl>  <int>  <dbl>    <dbl>  <int>
## 1 Fair      1.05    0.516   1610   64.0    3.64   1610    59.1    3.95   1610
## 2 Good      0.849    0.454   4906   62.4    2.17   4906    58.7    2.85   4906
## 3 Very ...  0.806    0.459  12082   61.8    1.38  12082    58.0    2.12  12082
## 4 Premi...  0.892    0.515  13791   61.3    1.16  13791    58.7    1.48  13791
## 5 Ideal     0.703    0.433  21551   61.7    0.719  21551    56.0    1.25  21551
## # ... with abbreviated variable names 1carat_mean, 2carat_sd, 3depth_mean,
## # 4depth_sd, 5table_mean, 6table_sd
```

Importantly, however, if a vector contains even one `NA`, remember the returned statistic for the entire vector will also be `NA`. The previous code will *NOT* return correct statistics if you have an `NA`.

You will want to remove them either by filtering rows for `complete.cases()` or omitting `NA`s from the summary statistic function. To illustrate, let's make some data missing in the data frame. Because there are 53,940 rows in the data frame, we can simply make some data in the last row missing. Let's change the data in `depth` and `table` variables to `NA` just to illustrate the point.

Hide

```
diamonds[53940, c("depth", "table")] # the current data
```

```
## # A tibble: 1 × 2
##   depth table
##   <dbl> <dbl>
## 1  62.2    55
```

Hide

```
diamonds[53940, c("depth", "table")] <- NA # set these cells to NA

tail(diamonds) # see missing values
```

```
## # A tibble: 6 × 10
##   carat cut      color clarity depth table price      x      y      z
##   <dbl> <ord>    <ord> <ord>   <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1  0.72 Premium D      SI1     62.7    59  2757  5.69  5.73  3.58
## 2  0.72 Ideal  D      SI1     60.8    57  2757  5.75  5.76  3.5
## 3  0.72 Good   D      SI1     63.1    55  2757  5.69  5.75  3.61
## 4  0.7 Very Good D      SI1     62.8    60  2757  5.66  5.68  3.56
## 5  0.86 Premium H      SI2     61     58  2757  6.15  6.12  3.74
## 6  0.75 Ideal  D      SI2     NA     NA  2757  5.83  5.87  3.64
```

Now filter out missing cases...

Hide

```
diamonds %>%
  filter(., complete.cases(.)) %>%
  group_by(., cut) %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
    .fns = list(mean = mean,
                 sd = sd,
                 n = length),
    .names = "{.col}_{.fn}"
  ))
```

```
## # A tibble: 5 × 10
##   cut      carat...1 carat...2 carat_n depth...3 depth...4 depth_n table...5 table...6 table_n
##   <ord>      <dbl>      <dbl>      <int>      <dbl>      <dbl>      <int>      <dbl>      <dbl>      <int>
## 1 Fair      1.05      0.516      1610      64.0      3.64      1610      59.1      3.95      1610
## 2 Good      0.849     0.454      4906      62.4      2.17      4906      58.7      2.85      4906
## 3 Very ...   0.806     0.459     12082      61.8      1.38     12082      58.0      2.12     12082
## 4 Premi...   0.892     0.515     13791      61.3      1.16     13791      58.7      1.48     13791
## 5 Ideal     0.703     0.433     21550      61.7      0.719     21550      56.0      1.25     21550
## # ... with abbreviated variable names 1carat_mean, 2carat_sd, 3depth_mean,
## # 4depth_sd, 5table_mean, 6table_sd
```

That's likely the easiest approach but if there is an NA in some variables and not others (e.g., only depth and table here), the variables with complete cases will also lose data. Note how n is the same for all variables.

What you may really want to do would be to compute the statistics by omitting NA s at the variable level. Now, across() has a way to handle this which involves passing an additional argument na.rm = TRUE which will adjust the functions, in this case mean and sd functions to include na.rm = TRUE . We can add this after .names . Unfortunately, however, length() annoyingly has no argument for removing NA s. If we drop it out of the list, we get the means and standard deviations.

Hide

```
diamonds %>%
  filter(., complete.cases(.)) %>%
  group_by(., cut) %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
    .fns = list(mean = mean,
                 sd = sd
               ),
    .names = "{.col}_{.fn}",
    na.rm = TRUE
  ))
```

```
## # A tibble: 5 × 7
##   cut      carat_mean carat_sd depth_mean depth_sd table_mean table_sd
##   <ord>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Fair      1.05      0.516      64.0      3.64      59.1      3.95
## 2 Good      0.849     0.454      62.4      2.17      58.7      2.85
## 3 Very Good   0.806     0.459      61.8      1.38      58.0      2.12
## 4 Premium     0.892     0.515      61.3      1.16      58.7      1.48
## 5 Ideal     0.703     0.433      61.7      0.719     56.0      1.25
```

So this doesn't solve the problem for all metrics you want to compute. What now?

Approach #1:

We can write our own function that contains NA removal in the same way. We just need to ensure we use na.rm and not something like remove.na . Let's name the function length_na .

Hide

```
length_na <- function(x, na.rm = FALSE) {
  # as length() substitute that calculates length with and without NAs
  if (na.rm) {
    x = length(na.omit(x))
  } else {
    x = length(x)
  }
  return(x)
}
```

And add `length_na` to the list...

Hide

```
diamonds %>%
  filter(., complete.cases(.)) %>%
  group_by(., cut) %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
    .fns = list(mean = mean,
                 sd = sd,
                 n = length_na
               ),
    .names = "{.col}_{.fn}",
    na.rm = TRUE
  ))
```

```
## # A tibble: 5 × 10
##   cut      carat...1 carat...2 carat_n depth...3 depth...4 depth_n table...5 table...6 table_n
##   <ord>    <dbl>    <dbl>    <int>    <dbl>    <dbl>    <int>    <dbl>    <dbl>    <int>
## 1 Fair      1.05     0.516    1610     64.0     3.64     1610     59.1     3.95     1610
## 2 Good      0.849     0.454    4906     62.4     2.17     4906     58.7     2.85     4906
## 3 Very ...  0.806     0.459   12082     61.8     1.38   12082     58.0     2.12   12082
## 4 Premi...  0.892     0.515   13791     61.3     1.16   13791     58.7     1.48   13791
## 5 Ideal     0.703     0.433   21550     61.7     0.719   21550     56.0     1.25   21550
## # ... with abbreviated variable names 1carat_mean, 2carat_sd, 3depth_mean,
## # 4depth_sd, 5table_mean, 6table_sd
```

But this removed NA row-wise and not for each column variable independently. The `carat` variable should not have the name number of cases as `depth` and `table`, both of which contain missing cases. This approach just did the same thing as filtering by complete cases as done earlier with `filter(., complete.cases(.))`.

Approach #2:

When functions do not contain argument for dealing with NA s, there is `na.omit()`, a function that takes an object and removes NA s. So you can just pass the variable to `na.omit()` and then wrap it in the metric function of interest. Also, because `na.rm = T` cannot be used for `length()`, `na.omit()` offers consistency across all functions and as a result, I believe, less confusion.

Unfortunately, accomplishing this task can be rather tricky and requires some new syntax. This requires usage of what's called a "lambda" technique. Using this type of syntax, we can pass functions to the `.fns` argument. The `?across()` documentation calls it "a purrr-style lambda" in the arguments section (for clarity, `purrr` is a

library). This approach can be a little bit confusing, so I'm going to show you an example, and then walk through it step by step.

We need to precede the function with `~` and reference the vector using `.x`. Let's do this and change the `.fns` argument slightly.

General Example:

```
name = ~function(na.omit(.x))
```

Hide

```
diamonds %>%
  group_by(., cut) %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
    .fns = list(mean = ~mean(na.omit(.x)),
                sd = ~sd(na.omit(.x)),
                n = ~length(na.omit(.x))),
    .names = "{.col}_{.fn}"
  )
)
```

```
## # A tibble: 5 × 10
##   cut      carat...1 carat...2 carat_n depth...3 depth...4 depth_n table...5 table...6 table_n
##   <ord>      <dbl>      <dbl>   <int>   <dbl>      <dbl>   <int>   <dbl>      <dbl>   <int>
## 1 Fair      1.05      0.516   1610    64.0      3.64    1610    59.1      3.95    1610
## 2 Good      0.849     0.454   4906    62.4      2.17    4906    58.7      2.85    4906
## 3 Very ...  0.806     0.459  12082    61.8      1.38   12082    58.0      2.12   12082
## 4 Premi...  0.892     0.515  13791    61.3      1.16   13791    58.7      1.48   13791
## 5 Ideal     0.703     0.433  21551    61.7      0.719  21550    56.0      1.25   21550
## # ... with abbreviated variable names 1carat_mean, 2carat_sd, 3depth_mean,
## # 4depth_sd, 5table_mean, 6table_sd
```

Great! Now `carat` contains one more case than the other variables. So what's the point of all of this? Well, you need to be careful not to apply functions and assume they are doing what you believe you are doing. You always need to be smarter than the code you use. Also, there is no single answer for dealing with data. Sometimes one approach will be appropriate and in other instances another approach will be. You as the data scientist need to know that there are different methods so that you can decide where to apply those different methods.

14.2 Grouping by multiple variables`

OK, now this gets exciting. When you want to group by additional variables, pass a new one to `group_by()`. Just to keep the output more simple, let's remove one summary function.

Hide


```

diamonds %>%
  group_by(., cut, clarity) %>%
  summarise(., across(.cols = c("carat", "depth", "table"),
    .fns = list(mean = ~mean(na.omit(.x)),
#              sd = ~sd(na.omit(.x)),
              n = ~length(na.omit(.x))),
    .names = "{.col}_{.fn}")
  )

```

```

## `summarise()` has grouped output by 'cut'. You can override using the `.groups`
## argument.

```

```

## # A tibble: 40 × 8
## # Groups:   cut [5]
##   cut    clarity carat_mean carat_n depth_mean depth_n table_mean table_n
##   <ord> <ord>         <dbl>   <int>    <dbl>   <int>    <dbl>   <int>
## 1 Fair   I1          1.36     210     65.7     210     58.1     210
## 2 Fair   SI2          1.20     466     64.4     466     58.8     466
## 3 Fair   SI1          0.965    408     63.9     408     59.1     408
## 4 Fair   VS2          0.885    261     63.6     261     59.1     261
## 5 Fair   VS1          0.880    170     62.9     170     60.4     170
## 6 Fair   VVS2          0.692     69     62.8     69     59.2     69
## 7 Fair   VVS1          0.665     17     60.4     17     61.2     17
## 8 Fair   IF           0.474      9     60.1      9     59.1      9
## 9 Good   I1           1.20      96     62.1     96     59.5     96
## 10 Good  SI2           1.04    1081     62.2    1081     58.9    1081
## # ... with 30 more rows

```

15 The Data Manipulation Workflow: Putting It All Together

Of course, all of this can be paired with `select()`, `mutate()`, `filter()`, etc. Here is the data manipulation workflow

```

dataframe %>% select(., ...) %>%      # select variables of interest
mutate(., ...) %>%                   # then create new variables
group_by(., ...) %>%                 # then group for subsetting
summarize(., ...)                    # then summarize

```